Module 6: Add Responsible AI Guardrails

Responsible AI Risk Checklist

Risk	Description	Guardrail
Privacy	Ticket text might contain personal details	Use only fictional tickets in the demo. Do not store or log real customer data.
Incorrect priority	AI might under- or over-prioritize a ticket	Always show the priority as a suggestion. Agents must confirm before acting.
Bias	Certain ticket types or language styles might be treated inconsistently	Review flagged results weekly. Update rules when patterns of unfair treatment appear.
Sensitive tickets	Security, accessibility, billing disputes, or angry messages need careful handling	Automatically flag for human review (see list below).
Over-automation	Agents might accept AI output without checking	Display a clear warning that all suggestions must be reviewed by a human before any customer response or routing action.

Cases That Must Always Be Reviewed by a Human

- Security incidents or possible unauthorized access
- Accessibility-related issues
- Billing disputes or refund requests
- Threatening, abusive, or highly escalated language
- Vague or unclear tickets where the AI confidence is low
- Any case the AI itself flags as uncertain

Privacy Note

This prototype uses only fictional sample tickets. No real customer data, personal information, or account details are collected or stored.

Warning Displayed in the Prototype

“These recommendations are AI-generated to speed up triage. Please review before routing or responding to the customer.”

Edge-Case Testing Notes

- Security ticket → flagged Yes (security incident)
- Accessibility ticket → flagged Yes (accessibility issue)
- Billing dispute → flagged Yes (billing dispute)

- Vague ticket (“This is not working”) → default to human review

This keeps the prototype helpful while ensuring humans remain responsible for important decisions.