Module 7: Add Agentic Workflow Thinking

Current Human + AI Workflow

1. Agent pastes a fictional ticket into the text box.
2. AI analyzes the text and returns a suggested summary, category, priority, recommended team, and human-review flag.
3. Agent reviews every suggestion.
4. If the human-review flag is Yes (or the agent disagrees with any output), the agent makes the final decision on category, priority, routing, and response.
5. Agent proceeds with the customer response or escalation. AI never sends a message or routes a ticket on its own.

Actions That Always Require Human Approval

- Final routing or assignment to a team
- Any customer-facing response
- Changing priority on a security, accessibility, or billing ticket
- Closing or escalating a ticket
- Updating the triage rules or prompt

Future Agentic Enhancement Plan

A later version could:

- Summarize the ticket and draft a suggested response automatically.
- Notify the recommended team when the human-review flag is No and the agent has approved the results.
- Update ticket status in a support system after the agent confirms.
- Still require explicit human approval before any notification, status change, or customer message is sent.

Think → Act → Observe Example

Think: The ticket says “I think someone accessed my account without permission.” This matches a security concern.

Act: Classify as Security Concern, set Priority to Urgent, recommend Security Team, and set Human Review to Yes.

Observe: The support agent reviews the flag, confirms the escalation, and assigns the ticket to the Security Team. The AI does not perform the assignment.

This workflow keeps the AI in an assisting role while humans retain control over important decisions.